

Project Milestone 2 Report

YAOJUN LIU, University of British Columbia, Canada
SASIVIMOL SIRIJANGKAPATTANA, University of British Columbia, Canada

1. Introduction

This Milestone 2 report documents progress on the first research question (RQ1) of the project. The objective is to examine whether repository- and user-level characteristics predict pull request (PR) resolution time.

The project investigates three research questions:

- **RQ1:** How do repository popularity (stars, forks) and user influence (followers, account age) relate to PR resolution time?
- **RQ2:** How do text characteristics of PRs (length and sentiment) vary across repositories of different popularity tiers and users of different influence levels?
- **RQ3:** How does AI coding agent contribution behavior vary across repositories with different popularity levels?

This milestone focuses exclusively on RQ1. RQ2 and RQ3 will be completed for the final report.

2. Data Wrangling and Feature Construction

Three datasets were used: pull requests, repositories, and users. All timestamps were converted to datetime format and merged using `repo_id` and `user_id`. Only closed PRs were retained to compute resolution time.

2.1 Resolution Time

PR resolution time was computed as:

$$resolution_time_hours = \frac{closed_at - created_at}{3600}$$

User account age was defined as:

$$user_account_age_days = created_at(PR) - created_at_user.$$

2.2 Zero-Inflated Predictors

Repository popularity (stars, forks) and user influence (followers) are heavily zero-inflated (60–80% zeros). To preserve dataset representativeness, zero values were retained. All predictors and the outcome were log-transformed using $\log(1 + x)$ to reduce skew.

3. Preliminary Results for RQ1

3.1 Correlation Analysis

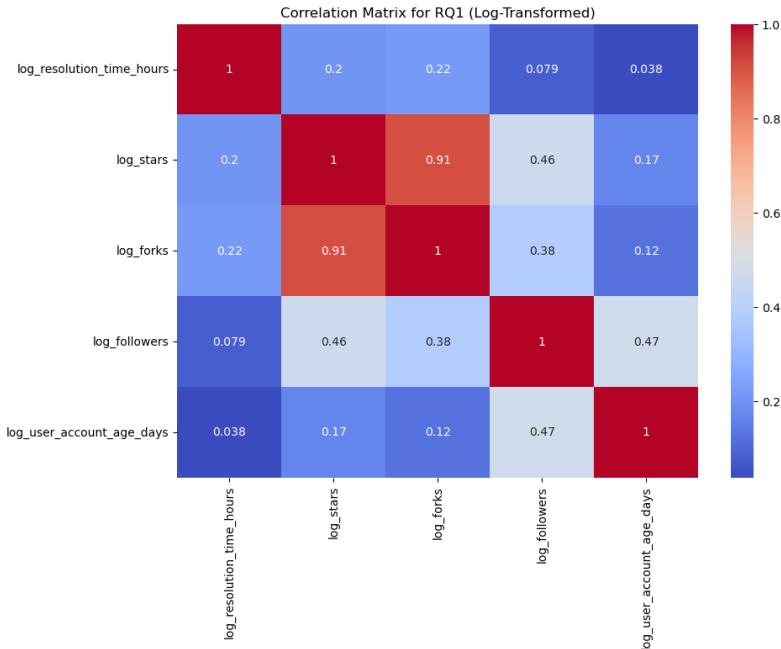
A correlation matrix of log-transformed variables shows weak associations between predictor variables and PR resolution time:

- $\text{corr}(\log(\text{resolution time}), \log(\text{stars})) \approx 0.20$
- $\text{corr}(\log(\text{resolution time}), \log(\text{forks})) \approx 0.22$
- $\text{corr}(\log(\text{resolution time}), \log(\text{followers})) \approx 0.079$

Authors' Contact Information: Yaojun Liu, University of British Columbia, Department of Computer Science, Okanagan, BC, Canada; Sasivimol Sirijangkappattana, University of British Columbia, Department of Computer Science, Okanagan, BC, Canada.

- 50 • $\text{corr}(\log(\text{resolution time}), \log(\text{user age})) \approx 0.038$

51 Details can be found in Figure 1.



75 Fig. 1. Pearson Correlation Matrix of Repository and User Metadata vs. Resolution Time.
76

77 All correlations are below 0.25, indicating negligible linear relationships between popularity/influence and PR processing time.
78
79

80 3.2 Regression Analysis 81

82 A log-linear regression model was fitted using 817,668 closed pull requests:
83

$$84 \quad \log(1+Y) = \beta_0 + \beta_1 \log(1+\text{stars}) + \beta_2 \log(1+\text{forks}) + \beta_3 \log(1+\text{followers}) + \beta_4 \log(1+\text{user age}) + \varepsilon. \\ 85$$

86 Model performance:

$$87 \quad R^2 = 0.049, \quad \text{Adjusted } R^2 = 0.049.$$

88 Thus, only about 4.9% of the variability in PR resolution time is explained by these predictors.
89

90 **Estimated coefficients (all $p < 0.001$):**
91

- 92 • β_1 (log-stars): +0.0163
- 93 • β_2 (log-forks): +0.2223
- 94 • β_3 (log-followers): -0.0120
- 95 • β_4 (log-account age): +0.0099

96 Although statistically significant due to the large sample size, these effects are extremely small
97 in magnitude.
98

Table 1. OLS Regression Results for Log Resolution Time

Variable	Coef	Std Err	t	P> t
const	0.1922	0.004	43.712	0.000
log_stars	0.0163	0.002	7.985	0.000
log_forks	0.2223	0.003	77.454	0.000
log_followers	-0.0120	0.001	-13.199	0.000
log_user_account_age_days	0.0099	0.001	14.716	0.000

3.3 Interpretation

These findings indicate that:

- Repository popularity has little to no predictive power for PR resolution time.
- User influence (followers, account age) also has negligible predictive impact.
- The small coefficient sizes and low R^2 suggest that PR resolution time is primarily driven by unobserved factors such as maintainer availability, CI outcomes, review workload, and project-specific processes.

This aligns with prior software engineering studies reporting that PR review speed is difficult to predict from repository- or user-level metadata alone.

4. Limitations

- Resolution time is highly skewed, with many short-duration PRs and a long tail of very long ones.
- Important complexity indicators (changed files, additions, deletions, comments) were not used in this milestone but may explain additional variance.
- Zero-inflation limits the predictive strength of popularity and influence metrics.

5. Next Steps: RQ2 and RQ3

RQ2: How do text characteristics of pull requests vary across repositories of different popularity levels and users of different influence levels?

We will analyze whether PR text features—such as sentiment and text length—differ across repositories with varying popularity (low, medium, high) and across users with different influence levels (based on follower count). This analysis will help determine whether more popular repositories or high-influence users tend to produce PRs with distinct linguistic or stylistic characteristics.

Methodology. PR, repository, and user data are merged using repo_id and user_id. Text-based features are computed by concatenating PR titles and bodies into a single field, from which text_length and sentiment polarity are derived using the TextBlob library. Repository popularity tiers are generated using star count quantiles, while user influence tiers are based on follower quantiles. Exploratory visualizations (e.g., boxplots) will then be used to compare sentiment and text length across popularity and influence levels.

Features Used. The analysis will rely on the following fields:

- title, body (raw PR text)
- full_text (combined text used for NLP features)
- text_length (character length of PR text)
- sentiment (sentiment polarity extracted via TextBlob)

- 148 • stars (repository popularity)
- 149 • followers (user influence)
- 150 • repo_id, user_id (merge keys)

151 152 **RQ3: How do AI agents behave across repositories of different popularity levels?**

153 We will analyze how AI-generated pull requests differ when submitted to repositories with varying
 154 popularity levels. The focus is on understanding whether AI agents contribute differently in terms
 155 of submission volume, merge outcomes, and PR lifetime across low, medium, and high-popularity
 156 repositories.

157 *Methodology.* PR metadata will be merged with repository information using repo_id. Repository
 158 popularity will be categorized into three tiers (low, medium, high) based on star count quantiles. For
 159 all PRs authored by AI agents, we will compute behavioral metrics—AI PR volume, closure/merge
 160 rate, and PR lifetime (closed_at – created_at)—within each popularity tier. Exploratory data analysis,
 161 including summary statistics and visualizations, will be used to compare AI behavior across these
 162 tiers.

163 *Features Used.* The analysis will rely on the following fields:

- 164 • agent (AI-agent: Claude_Code, Copilot,...)
- 165 • repo_id (repository identifier)
- 166 • stars (repository popularity measure)
- 167 • state (merged/closed/open)
- 168 • created_at and closed_at (timestamps used to compute PR lifetime)
- 169 • id (pull request identifier)

170 171 **6. GitHub Repository**

172 Code and data processing pipeline:

173 <https://github.com/liuyaojun1/542.git>

174

175

176

177

178

179

180

181

182

183

184

185

186

187

188

189

190

191

192

193

194

195

196